

Data Visualisation and Manipulation

DATA VISUALISATION

In this document we are analysing the covid dataset. In order to do so the first step was to require ggplot2, and dplyr libraries.

```
library(readr)
library(ggplot2)
library(dplyr)
library(tidyr)
library(magrittr)
library(knitr)
library(MASS)
```

The second step was to upload the dataset, and to eliminate the index column that was not needed

```
covid <- read_csv("covid.csv")
data<- covid[,-1]
View(data)
```

Then, we inspected the dataset before visualizing it in order to understand what it was about.

```
data%>% glimpse
range(data$year)
data%>% distinct(year)
length(unique(data$country))
length(unique(data$continent))
data%>% distinct(year, cases)
sum(is.na(data$deaths))
sum(is.na(data$cases))
```

From the above code it emerged that: the data set is made of 61900 observations and 12 variables. The year inspected is from 31-12-2019 to 14-12-2020. There are 214 countries from 6 continents. We also checked that there were no missing values

DATA VISUALISATION WITH GGPLO2

For data visualization the ggplot2 library was used, sometimes together with some dplyr functions. The first plot compares the distribution of COVID cases by country, in this case we chose as example Belgium and Spain.


```
data%>%
  filter(country %in% c("Belgium", "Spain")) %>%
  ggplot(
```

```

aes(x = cases_cum, y = ..density.., color = I("white"), fill = country)
) +
geom_histogram(alpha = 0.6, position = "identity") +
scale_fill_manual(values = c("grey60", "darkorange")) +
theme_bw()

```

'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.




The same plot was then used to show the distribution of deaths for the same country. Note that the filter() function is part of the dplyr packaged and is used to extract the countries of interest

```

data %>%
  filter(country %in% c("Belgium", "Spain"))%>%
  ggplot(
    aes(x= deaths_cum, y= ..density.., color=I("white"), fill=country)
  )+
  geom_histogram(alpha=0.6, position="identity")+
  scale_fill_manual(values=c("grey60", "darkorange"))+
  theme_bw()

```


'stat_bin()' using 'bins = 30'. Pick better value with 'binwidth'.



BOXPLOT


The following plot shows the distribution of cases by month

```
data %>%
  ggplot(aes(x = month, y = cases, fill = I("skyblue2"))) +
  geom_boxplot(position = "dodge") +
  labs(x = "") +
  theme_bw() +
  theme(axis.text.x = element_text(angle = 90, vjust = .4, size = 10))
```




Same plot can be used to show the distribution of deaths.

```
data%>%
  ggplot(aes(x=month, y=deaths, fill=I("blue")))+
  geom_boxplot(position="dodge")+
  labs(x="")+
  theme_bw()+
  theme(axis.text.x=element_text(angle=90, vjust=.4, size=10))
```



Then, we can visualise changes in the number of cases in each continent on different months

```
ggplot(data, aes(month, cases))+
  geom_line(aes(group= continent), colour="grey50")+
  geom_point(aes(colour=continent))
```



Then we decided to summarise the statistics. In this case to inspect the number of deaths

```
by(data$cases, data$deaths, summary)
by(data$cases, data$deaths, sd)

data %>%
  group_by(deaths) %>%
  summarise(
    Minimum = min(cases),
    Q1 = quantile(cases, probs = .25),
    Median = quantile(cases, probs = .5),
    Mean = mean(cases),
    Q3 = quantile(cases, probs = .75),
    Maximum = max(cases),
    SD = sd(cases),
    CV = abs(Mean)/SD,
    IQR = Q3 - Q1
  )
```

GROUPED BOXPLOT


With the filter function we created a time series to show the distribution of cases in each continent

```
data%>%
  filter(
    month %in% c("01", "02", "03", "04", "05", "06", "07", "08", "09", "10", "11", "12"),
```

```

  grepl("[aeiouy]$", country)
) %>%
ggplot(aes(x = continent, y = cases, fill = month)) +
  geom_boxplot(position = "dodge") +
  theme(legend.position = "top", axis.text.x = element_text(angle = 90))

```




A time series was then created for the country Belgium.

```

data %>% filter(country == "Belgium") %>%
  ggplot(aes(x = date, y = cases, color = date)) + geom_point() + geom_line() +
  scale_y_continuous(labels = scales::percent) +
  theme_bw() +
  theme(legend.position = "bottom")

```



DATA MANUPULATION WITH DPLYR

Now we perform some data manipulation with the dplyr package. Suppose that we want to create a new dataframe with some of the variables from our original data. We use the `select()` function, and name the new dataset “data2”. We added `dplyr::select` in order to avoid conflicts with the MASS library.

```
data2<- data%>%
  dplyr::select(deaths)
data2
```

```
## # A tibble: 61,900 x 1
##   deaths
##   <dbl>
## 1      0
## 2      0
## 3      0
## 4      0
## 5      0
## 6      0
## 7      0
## 8      0
## 9      0
## 10     0
## # ... with 61,890 more rows
```

Now suppose we want to rename the column `deat` with the new name being “KO”. We used for this purpose the `rename()` function.


```
data3<- rename(data2, KO= deaths)
data3
```

```
## # A tibble: 61,900 x 1
##       KO
##   <dbl>
## 1     0
## 2     0
## 3     0
## 4     0
## 5     0
## 6     0
## 7     0
## 8     0
## 9     0
## 10    0
## # ... with 61,890 more rows
```

Then we used the filter() function in order to choose a subset of the original data, and retain only the values in which the variable death is equal to ten. Then we added the number “30” to retain also the cases in which the number of deaths was 30.

```
data4<-filter(data, deaths == "10")
data4
```

```
## # A tibble: 488 x 12
##   date      day month year cases deaths country  code population continent
##   <date>    <chr> <chr> <dbl> <dbl> <dbl> <chr>    <chr>    <dbl> <chr>
## 1 2020-04-27 27   04   2020    68    10 Afghanistan AF    38041757 Asia
## 2 2020-07-27 27   07   2020   106    10 Afghanistan AF    38041757 Asia
## 3 2020-08-13 13   08   2020    76    10 Afghanistan AF    38041757 Asia
## 4 2020-08-25 25   08   2020    71    10 Afghanistan AF    38041757 Asia
## 5 2020-09-16 16   09   2020    40    10 Afghanistan AF    38041757 Asia
## 6 2020-11-14 14   11   2020    66    10 Afghanistan AF    38041757 Asia
## 7 2020-12-06 06   12   2020   234    10 Afghanistan AF    38041757 Asia
## 8 2020-12-11 11   12   2020    63    10 Afghanistan AF    38041757 Asia
## 9 2020-11-27 27   11   2020   656    10 Albania    AL    2862427 Europe
## 10 2020-04-23 23   04   2020    99    10 Algeria    DZ    43053054 Africa
## # ... with 478 more rows, and 2 more variables: cases_cum <dbl>,
## #   deaths_cum <dbl>
```

```
data5<- filter(data, deaths %in% c("10", "30"))
data5
```

```
## # A tibble: 604 x 12
##   date      day month year cases deaths country  code population continent
##   <date>    <chr> <chr> <dbl> <dbl> <dbl> <chr>    <chr>    <dbl> <chr>
## 1 2020-04-27 27   04   2020    68    10 Afghanistan AF    38041757 Asia
## 2 2020-06-08 08   06   2020   791    30 Afghanistan AF    38041757 Asia
## 3 2020-07-27 27   07   2020   106    10 Afghanistan AF    38041757 Asia
## 4 2020-08-13 13   08   2020    76    10 Afghanistan AF    38041757 Asia
## 5 2020-08-25 25   08   2020    71    10 Afghanistan AF    38041757 Asia
```

```
## 6 2020-09-16 16 09 2020 40 10 Afghanistan AF 38041757 Asia
## 7 2020-11-14 14 11 2020 66 10 Afghanistan AF 38041757 Asia
## 8 2020-12-06 06 12 2020 234 10 Afghanistan AF 38041757 Asia
## 9 2020-12-11 11 12 2020 63 10 Afghanistan AF 38041757 Asia
## 10 2020-11-27 27 11 2020 656 10 Albania AL 2862427 Europe
## # ... with 594 more rows, and 2 more variables: cases_cum <dbl>,
## # deaths_cum <dbl>
```

We can use the function `summarise()` and take the mean and the median of the variable death.

```
summarise(data2, d_mean= mean(deaths), d_med= median(deaths))
```

```
## # A tibble: 1 x 2
##   d_mean d_med
##   <dbl> <dbl>
## 1  26.1    0
```

We can use the function `group by()` and `summarise at()` to create another dataframe with the selected variables

```
a<- data%>%group_by(deaths)%>%
  summarise_at(vars(cases, population), funs(n(), mean(., na.rm=TRUE)))
```

```
## Warning: 'funs()' was deprecated in dplyr 0.8.0.
## Please use a list of either functions or lambdas:
##
##   # Simple named list:
##   list(mean = mean, median = median)
##
##   # Auto named with 'tibble::lst()':
##   tibble::lst(mean, median)
##
##   # Using lambdas
##   list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
## This warning is displayed once every 8 hours.
## Call 'lifecycle::last_warnings()' to see where this warning was generated.
```

```
a
```

```
## # A tibble: 1,049 x 5
##   deaths cases_n population_n cases_mean population_mean
##   <dbl> <int> <int> <dbl> <dbl>
## 1 -1918 1 1 -372 46937060
## 2 -443 1 1 237 6415851
## 3 -31 1 1 577 60359546
## 4 -12 1 1 218 1798506
## 5 -5 1 1 466 4904240
## 6 -3 1 1 75 10650000
## 7 -2 1 1 3172 46937060
## 8 -1 1 1 121 10650000
## 9 0 36728 36728 26.0 25743880.
## 10 1 4464 4464 124. 29333910.
## # ... with 1,039 more rows
```

The `min_rank()` function is a function that returns the same values as `rank` when the `ties_method` is set to “min”, that is, ties are assigned the minimum ranking possible

```
rank<- data %>% group_by(deaths)%>%filter(min_rank(desc(population))==50)%>%
  dplyr::select(deaths, year, population)
rank
```

```
## # A tibble: 56 x 3
## # Groups:   deaths [20]
##   deaths year population
##   <dbl> <dbl>      <dbl>
## 1     20  2020    44780675
## 2     38  2020    10047719
## 3     38  2020    10047719
## 4     38  2020    10047719
## 5     47  2020    11455519
## 6     45  2020    11455519
## 7     45  2020    11455519
## 8     31  2020    18952035
## 9     42  2020    18952035
## 10    42  2020    18952035
## # ... with 46 more rows
```

Finally, we can use the `join()` function. Specifically, `inner_join()` returns rows when there is a match in both tables. In this case, I am merging `rank` and `data4` with `deaths` as primary key

```
d<- inner_join(rank,data4, by="deaths")
d
```

```
## # A tibble: 1,952 x 14
## # Groups:   deaths [1]
##   deaths year.x population.x date      day month year.y cases country code
##   <dbl> <dbl>      <dbl> <date>   <chr> <chr> <dbl> <dbl> <chr> <chr>
## 1     10  2020    100388076 2020-04-27 27    04    2020    68 Afghani~ AF
## 2     10  2020    100388076 2020-07-27 27    07    2020   106 Afghani~ AF
## 3     10  2020    100388076 2020-08-13 13    08    2020    76 Afghani~ AF
## 4     10  2020    100388076 2020-08-25 25    08    2020    71 Afghani~ AF
## 5     10  2020    100388076 2020-09-16 16    09    2020    40 Afghani~ AF
## 6     10  2020    100388076 2020-11-14 14    11    2020    66 Afghani~ AF
## 7     10  2020    100388076 2020-12-06 06    12    2020   234 Afghani~ AF
## 8     10  2020    100388076 2020-12-11 11    12    2020    63 Afghani~ AF
## 9     10  2020    100388076 2020-11-27 27    11    2020   656 Albania AL
## 10    10  2020    100388076 2020-04-23 23    04    2020    99 Algeria DZ
## # ... with 1,942 more rows, and 4 more variables: population.y <dbl>,
## # continent <chr>, cases_cum <dbl>, deaths_cum <dbl>
```